Practical aspects of deep learning

10/10 points (100%)

Teste, 10 questions

✔ Parabéns! Você foi aprovado!

Próximo item

✔ 1/1 pontos

1.
If you have 10,000,000 examples, how would you split the train/dev/test set?

- ◉ 98% train . 1% dev . 1% test

Correto

- ○ 60% train . 20% dev . 20% test
- ○ 33% train . 33% dev . 33% test

✔ 1/1 pontos

2.
The dev and test set should:

- ◉ Come from the same distribution

Correto

- ○ Come from different distributions
- ○ Be identical to each other (same (x,y) pairs)

Have the same number of examples

Practical aspects of deep learning

10/10 points (100%)

Teste, 10 questions

1/1 pontos

3.

If your Neural Network model seems to have high variance, what of the following would be promising things to try?

- [x] Add regularization

Correto

- [x] Get more training data

Correto

- [ ] Increase the number of units in each hidden layer

Não selecionado está correto

- [ ] Get more test data

Não selecionado está correto

- [ ] Make the Neural Network deeper

Não selecionado está correto

1/1 pontos

4.

You are working on an automated check-out kiosk for a supermarket, and are building a classifier for apples, bananas and oranges. Suppose your classifier obtains a training set error of 0.5%, and a dev set error of 7%. Which of the following are promising things to try to improve your classifier? (Check all that apply.)

- [x] Increase the regularization parameter lambda

Correto

- [ ] Decrease the regularization parameter lambda

Não selecionado está correto

- [x] Get more training data

Correto

- [ ] Use a bigger neural network

Não selecionado está correto

1/1 pontos

5.
What is weight decay?

- (•) A regularization technique (such as L2 regularization) that results in gradient descent shrinking the weights on every iteration.

Correto

- ( ) Gradual corruption of the weights in the neural network if it is trained on noisy data.
- ( ) The process of gradually decreasing the learning rate during training.
- ( ) A technique to avoid vanishing gradient by imposing a ceiling on the values of the weights.

1/1 pontos

Practical aspects of deep learning

10/10 points (100%)

Teste, 10 questions

6.
What happens when you increase the regularization hyperparameter lambda?

- [x] Weights are pushed toward becoming smaller (closer to 0)

Correto

- [ ] Weights are pushed toward becoming bigger (further from 0)
- [ ] Doubling lambda should roughly result in doubling the weights
- [ ] Gradient descent taking bigger steps with each iteration (proportional to lambda)

1/1 pontos

7.
With the inverted dropout technique, at test time:

- [ ] You do not apply dropout (do not randomly eliminate units), but keep the 1/keep_prob factor in the calculations used in training.
- [ ] You apply dropout (randomly eliminating units) but keep the 1/keep_prob factor in the calculations used in training.
- [ ] You apply dropout (randomly eliminating units) and do not keep the 1/keep_prob factor in the calculations used in training
- [x] You do not apply dropout (do not randomly eliminate units) and do not keep the 1/keep_prob factor in the calculations used in training

Correto

1/1 pontos

8.
Increasing the parameter keep_prob from (say) 0.5 to 0.6 will likely cause the following: (Check the two that apply)

- [ ]

- [ ] Increasing the regularization effect

Não selecionado está correto

- [ ] Reducing the regularization effect

Correto

- [ ] Causing the neural network to end up with a higher training set error

Não selecionado está correto

- [ ] Causing the neural network to end up with a lower training set error

Correto

1/1 pontos

9.
Which of these techniques are useful for reducing variance (reducing overfitting)? (Check all that apply.)

- [ ] Gradient Checking

Não selecionado está correto

- [ ] Vanishing gradient

Não selecionado está correto

- [ ] Exploding gradient

Não selecionado está correto

- [ ] Dropout

Correto

Practical aspects of deep learning

10/10 points (100%)

Teste, 10 questions

- [x] L2 regularization

Correto

- [x] Xavier initialization

Não selecionado está correto

- [x] Data augmentation

Correto

1/1 pontos

10.

Why do we normalize the inputs x?

- ( ) It makes the parameter initialization faster
- ( ) It makes it easier to visualize the data
- (•) It makes the cost function faster to optimize

Correto

- ( ) Normalization is another word for regularization--It helps to reduce variance